

# notebook

January 19, 2026

```
[64]: import pandas as pd  
import matplotlib.pyplot as plt  
import matplotlib.colors as mcolors
```

```
[65]: df = pd.read_csv("data/Canadian_Health_Survey_Sample_cleaned.csv")
```

## 0.0.1 Analyzing dataset

```
[66]: df.head()
```

```
[66]:
```

	Province	Gender	Age	Income	BMI	PhysicalActivity	\
0	Ontario	Female	59	80572.0	25.7		3.8
1	New Brunswick	Male	38	125739.0	24.2		3.4
2	Saskatchewan	Male	30	75947.0	28.3		3.9
3	Saskatchewan	Female	79	113966.0	31.2		1.6
4	Newfoundland and Labrador	Female	24	101828.0	27.8		4.5

	Smoking	SelfRatedHealth	StressLevel	BMI_Category	IncomeBracket	\
0	Yes	Good	6	Overweight	80-110k	
1	No	Good	9	Normal	>110k	
2	No	Excellent	10	Overweight	50-80k	
3	No	Fair	6	Obese	>110k	
4	Yes	Good	5	Overweight	80-110k	

	SelfRatedHealth_Num
0	3
1	3
2	5
3	2
4	3

```
[67]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 825 entries, 0 to 824  
Data columns (total 12 columns):  
 #   Column           Non-Null Count  Dtype     
 ---  --     
 1   ---  -----  -----  -----
```

```

0 Province          825 non-null   object
1 Gender            825 non-null   object
2 Age               825 non-null   int64
3 Income            825 non-null   float64
4 BMI               825 non-null   float64
5 PhysicalActivity 825 non-null   float64
6 Smoking           825 non-null   object
7 SelfRatedHealth  825 non-null   object
8 StressLevel       825 non-null   int64
9 BMI_Catagory     825 non-null   object
10 IncomeBracket    825 non-null   object
11 SelfRatedHealth_Num 825 non-null   int64
dtypes: float64(3), int64(3), object(6)
memory usage: 77.5+ KB

```

```
[68]: df.describe(include='all')
```

	Province	Gender	Age	Income	BMI	\
count	825	825	825.000000	825.000000	825.000000	
unique	10	2	NaN	NaN	NaN	
top	Quebec	Male	NaN	NaN	NaN	
freq	97	426	NaN	NaN	NaN	
mean	NaN	NaN	49.408485	77545.210909	27.472039	
std	NaN	NaN	18.187646	22145.238035	4.953750	
min	NaN	NaN	18.000000	5000.000000	13.500000	
25%	NaN	NaN	33.000000	64115.000000	24.200000	
50%	NaN	NaN	49.000000	78130.000000	27.488861	
75%	NaN	NaN	66.000000	90100.000000	30.700000	
max	NaN	NaN	80.000000	250000.000000	60.000000	
	PhysicalActivity	Smoking	SelfRatedHealth	StressLevel	BMI_Catagory	\
count	825.000000	825	825	825.000000	825	
unique	NaN	2	5	NaN	4	
top	NaN	No	Good	NaN	Overweight	
freq	NaN	640	254	NaN	333	
mean	3.503246	NaN	NaN	5.294545	NaN	
std	1.899122	NaN	NaN	2.121156	NaN	
min	0.000000	NaN	NaN	1.000000	NaN	
25%	2.100000	NaN	NaN	4.000000	NaN	
50%	3.500000	NaN	NaN	5.000000	NaN	
75%	4.700000	NaN	NaN	7.000000	NaN	
max	10.400000	NaN	NaN	10.000000	NaN	
	IncomeBracket	SelfRatedHealth_Num				
count	825	825.000000				
unique	4	NaN				
top	50-80k	NaN				

freq	394	NaN
mean	NaN	3.208485
std	NaN	1.147042
min	NaN	1.000000
25%	NaN	2.000000
50%	NaN	3.000000
75%	NaN	4.000000
max	NaN	5.000000

## 0.0.2 Cleaning & Standardizing Categories

- Normalizing gender values and smoking values

```
[69]: # normalizing gender values to Male or Female
df['Gender'] = df['Gender'].str.lower()
df['Gender'] = df['Gender'].replace({
    'male': 'Male',
    'm': 'Male',
    'female': 'Female',
    'f': 'Female'
})

# normalizing smoking to Yes or No
df['Smoking'] = df['Smoking'].str.lower()
df['Smoking'] = df['Smoking'].replace({
    'yes': 'Yes',
    'y': 'Yes',
    'no': 'No',
    'n': 'No'
})
```

## 0.0.3 Data Cleaning

- Filled missing BMI and Income values using averages (mean/median) to preserve dataset size without distorting quantitative distributions.
- Dropped rows with missing StressLevel because it is ordinal, and imputing an ordered scale would distort the meaning of the categories.
- Filled missing PhysicalActivity values using the mean, since it is a discrete numeric variable and will not affect categorical encodings in later visualizations.

```
[70]: df['BMI'] = df['BMI'].fillna(df['BMI'].mean())
df['Income'] = df['Income'].fillna(df['Income'].median())
df = df.dropna(subset=['StressLevel'])
df['PhysicalActivity'] = df['PhysicalActivity'].fillna(df['PhysicalActivity'].mean())
```

## 0.1 Enforcing Correct Data Types

```
[71]: df['Age'] = df['Age'].astype(int)
df['Income'] = df['Income'].astype(float)
df['BMI'] = df['BMI'].astype(float)
df['StressLevel'] = df['StressLevel'].astype(int)
```

### 0.1.1 Enhancing the Dataset with Derived Features

```
[72]: # create new column
df['BMI_Category'] = None

# create column values based on constraints
df.loc[df['BMI'] < 18.5, 'BMI_Category'] = 'Underweight'
df.loc[(df['BMI'] >= 18.5) & (df['BMI'] <= 24.9), 'BMI_Category'] = 'Normal'
df.loc[(df['BMI'] >= 25) & (df['BMI'] <= 29.9), 'BMI_Category'] = 'Overweight'
df.loc[df['BMI'] >= 30, 'BMI_Category'] = 'Obese'

# create new column
df['IncomeBracket'] = None

# create column values based on constraints
df.loc[df['Income'] < 50000, 'IncomeBracket'] = '<50k'
df.loc[(df['Income'] >= 50000) & (df['Income'] < 80000), 'IncomeBracket'] = '50-80k'
df.loc[(df['Income'] >= 80000) & (df['Income'] < 110000), 'IncomeBracket'] = '80-110k'
df.loc[df['Income'] >= 110000, 'IncomeBracket'] = '>110k'
```

### 0.1.2 Bar Chart - Average Life/Health Proxy by Province Analysis

- This bar chart is created to compare the average self-rated health score across provinces.

```
[73]: # create new column
df['SelfRatedHealth_Num'] = None

# create column values based on constraints
df.loc[df['SelfRatedHealth'] == 'Poor', 'SelfRatedHealth_Num'] = 1
df.loc[df['SelfRatedHealth'] == 'Fair', 'SelfRatedHealth_Num'] = 2
df.loc[df['SelfRatedHealth'] == 'Good', 'SelfRatedHealth_Num'] = 3
df.loc[df['SelfRatedHealth'] == 'Very Good', 'SelfRatedHealth_Num'] = 4
df.loc[df['SelfRatedHealth'] == 'Excellent', 'SelfRatedHealth_Num'] = 5

# group each province
groups = df.groupby('Province')

# get the average score for each province
avg_score = groups['SelfRatedHealth_Num'].mean()
```

```

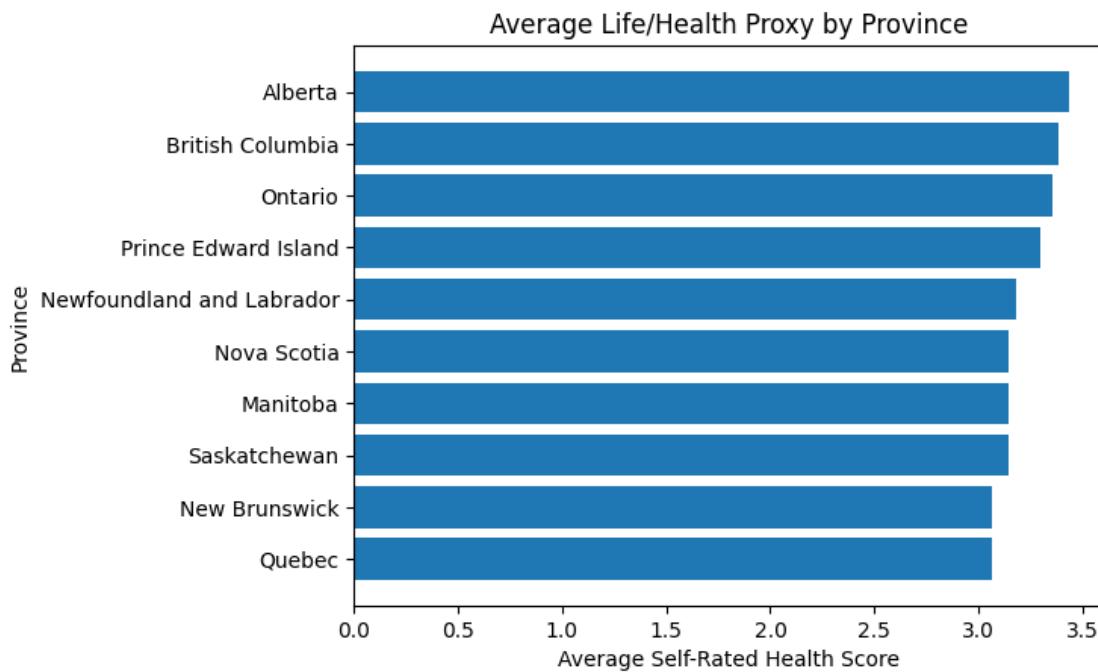
avg_score = avg_score.sort_values() # better visualization

# plot
plt.barh(avg_score.index, avg_score.values)

plt.xlabel('Average Self-Rated Health Score')
plt.ylabel('Province')
plt.title('Average Life/Health Proxy by Province')

plt.show()

```



### 0.1.3 Scatter Plot - Income vs BMI Analysis

- This scatter plot is created to examine the relationship between Income and BMI while also showing how StressLevel and Gender influence the pattern.

```

[74]: # create two new dataframes only containing the gender column where one is only
      ↵males and the other is only females
male_df = df[df['Gender'] == 'Male']
female_df = df[df['Gender'] == 'Female']

# for males
plt.scatter(
    male_df['Income'],
    male_df['BMI'],

```

```

c = male_df['StressLevel'], # colormap with stress level
cmap = 'viridis',
marker = 'o', # circle for males
label = 'Male' # for legend
)

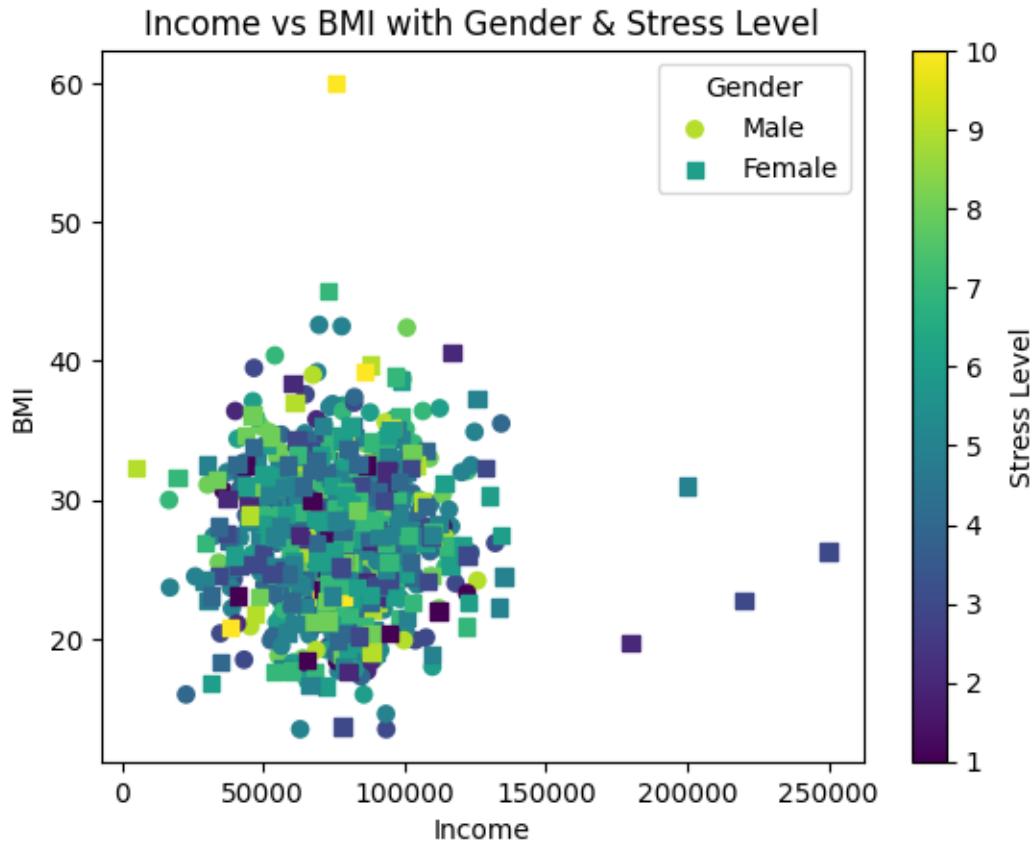
# for females
plt.scatter(
    female_df['Income'],
    female_df['BMI'],
    c = female_df['StressLevel'], # colormap with stress level
    cmap = 'viridis',
    marker = 's', # squares for females
    label = 'Female' # for legend
)

# Colorbar for StressLevel
colorbar = plt.colorbar()
colorbar.set_label('Stress Level')

plt.xlabel('Income')
plt.ylabel('BMI')
plt.title('Income vs BMI with Gender & Stress Level')

plt.legend(title='Gender')
plt.show()

```



#### 0.1.4 Diverging Bar Chart - Life Proxy Difference by Province Analysis

- This diverging bar chart is used to show how each province's average self-rated health score differs from the overall national average.
- The values can be both above and below zero, a diverging visualization is the correct choice as it highlights positive vs. negative deviation around a meaningful midpoint.

```
[75]: groups = df.groupby('Province')

# get the average score for each province
province_avg = groups['SelfRatedHealth_Num'].mean()

# get the overall average
ovr_avg = df['SelfRatedHealth_Num'].mean()

life_proxy = province_avg - ovr_avg
life_proxy = life_proxy.sort_values()
life_proxy = life_proxy.astype(float)

# color map
```

```

cmap = plt.cm.coolwarm # choose coolwarm color map

low = life_proxy.min()
high = life_proxy.max()
center = 0

color_map = mcolors.TwoSlopeNorm(vmin = low, vcenter = center, vmax = high)
colors = cmap(color_map(life_proxy.values)) # mapping each value to a color

plt.barh(life_proxy.index, life_proxy.values, color = colors)

plt.xlabel('Δ Life Proxy')
plt.ylabel('Province')
plt.title('Δ Life Proxy by Province')

plt.show()

```

